
PROFESSIONAL SUMMARY

Experienced Data Analyst with nearly 4 years of experience who focused on business outcomes with SQL/Python, Power BI/DAX, Excel and Azure. At DRG, built ADF pipelines that cut daily refresh ~50% and automated ~16 staff-hours/week. Delivered 15+ executive dashboards enabling self-serve analytics and faster decisions across Ops & Finance. Databricks (DLT/Delta) personal projects demonstrate end-to-end data ops tied to clear KPIs. Translate ambiguous asks into SQL/Python analyses and clear KPIs/OKRs, pairing A/B testing with data quality & governance to drive confident, SLA-aligned decisions.

SKILLS

- **Analytics & BI:** Power BI, DAX, Power Query (M), KPI design, data visualization, dashboarding, row-level security (RLS), data storytelling, self-serve BI, Excel (Pivot Tables, VBA Automation)
- **Programming & Data:** SQL (CTEs, window functions, performance tuning), Python (Pandas, NumPy, PySpark, SQLAlchemy, Flask), Excel (advanced formulas, PivotTables, VBA), REST APIs (JSON/XML)
- **Data Modeling & Warehousing:** Star/Snowflake schema, dimensional modeling, SCD1/SCD2, semantic models, ELT/ETL, data marts
- **Cloud Platforms:** Azure Data Factory (ADF), Databricks (Delta/Autoloader, DLT), ADLS Gen2, Azure SQL Database, PostgreSQL, Git/GitHub, CI/CD
- **Statistics & Experimentation:** Hypothesis testing, A/B testing, regression, cohort analysis, segmentation, time-series analysis, sample sizing
- **Quality & Governance:** Data profiling, validation rules, anomaly detection, documentation, lineage, UAT, SLA/OKR alignment
- **Ways of Working:** Stakeholder management, requirements gathering, JIRA, cross-functional collaboration

WORK HISTORY

Data Analyst**May 2024 – Present****The DRG, United Kingdom**

- Built a data & reporting stack using Azure, Python, SQL, Power Bi and Excel, enabling near-real-time insights for business stakeholders and faster decision-making.
- **Reduced daily refresh ~50% (120 to 60 mins)** across core ADF pipelines via SQL optimization, **incremental loads**, and reusable **Python** transforms data ready for **6am** Ops/Finance huddles.
- **Automated ~16 staff-hours/week** by orchestrating **VBA/Python** jobs and SQL scheduling to burst standardized PDF packs to **20+ sites**, cutting manual reporting and formatting defects.
- **Designed and analyzed A/B tests** (hypotheses, sample sizing, lift analysis, decision memos) to guide campaign rollout for Loyalty/Marketing.
- **Integrated REST/JSON and web data** into normalized relational models (**Azure PostgreSQL**); improved query ability and shortened time-to-insight for weekly performance reviews.
- Implemented automated Power BI dashboards for **nearly real-time** reporting and strategic decision-making, increasing quarterly revenue by 15% through enhanced sales insights.
- Automated daily REST API ingestion and web scraping (**BeautifulSoup, Selenium**); normalised **JSON/XML** to relational tables, improving database query performance and usability for analysts.

Data Analyst**April 2021 - March 2022****NSEIT, Nashik**

- **Streamlined SQL/Python ETL** for operations and examination data, **freeing ~10 hours/week** and improving report turnaround for internal stakeholders.
- Implemented data validation & reconciliation checks (row counts, referential integrity, anomaly flags) to raise reporting reliability and reduce rework.
- Produced and maintained **20+ decision-support reports/dashboards** with defined **KPIs** and **SLAs**, improved stakeholder satisfaction and consistency.
- **Standardized documentation and lineage** with a six-member data team (Git/GitHub), improving auditability and onboarding.

Data Analyst**March 2020 – March 2021****Swami Vivekanand Education and Research Centre, Mumbai**

- Built **Power BI** dashboards and optimized **SQL** models to surface customer analytics, **reducing data processing time ~20%**.
- Developed Power BI dashboards for real-time visualization of key metrics.

- Partnered with Marketing on **KPI definitions** and segmentation; **increased stakeholder engagement ~30%** through actionable visuals.
- Partnered with Admissions/Marketing on **KPI definitions** and segmentation; **increased stakeholder engagement ~30%** through actions

EDUCATION

- MSc in Advance Computer Science with Data science (2:1) University of Strathclyde, School of Computing Science **January 2023 – January 2024**
- Bachelor of Computer Applications (2:1) Savitribai Phule Pune university **May 2016 – March 2020**

PROJECTS

End-to-End Azure Data Engineering Lakehouse ([GitHub Link](#))

- Medallion (Bronze/Silver/Gold) on **ADLS Gen2**; parameterized, **metadata-driven ADF** with incremental loads, backfill controls, and JSON CDC watermarking.
- **Databricks Autoloader + Structured Streaming** refine Bronze to Silver **Delta**; **Unity Catalog governance** (metastore, external locations, access connector).
- **Delta Live Tables** for Gold: **SCD 2 dims** (sequence_by), SCD1 fact upserts, **Expectations + visual lineage DAG**; analytics-ready for **SQL Warehouse**/Power BI.
- Ops & DevEx: Python utilities (cleaning, de-dupe), checkpointing & idempotent runs, empty-file cleanup, **Logic Apps** alerts, secrets via **Managed Identity/Key Vault**, CI/CD with **Databricks Asset Bundles** + GitHub

DataBricks-Declarative-Pipelines ([GitHub Link](#))

- Built End-to-end **DLT Lakehouse** with **Medallion**; **unified batch + streaming** and automatic **DAG/lineage**.
- **Bronze**: streaming ingests with **Expectations** (warn/drop/fail) and Append Flow to merge multi-region sources; runtime audit metrics.
- Silver: **Auto-CDC** Type-1 upserts using reusable Python **utilities** (casting, de-dup, common transforms) for **incremental** accuracy.
- **Gold & Ops**: **SCD2** dimensions, Type-1 fact, and a **materialized revenue view**; **Unity Catalog governance**, parameterized dev/prod configs, **SQL seed scripts**, and **alerting**.

Azure Data Factory End-to-End (On-Prem + API + SQL) ([GitHub Link](#))

- ADF hybrid ingestion: **SHIR (on-prem shares)**, **HTTP/REST**, and Azure SQL **incremental loads** without watermark tables to **idempotent re-runs** and drift-tolerant landings.
- Orchestrated Bronze, Silver, Gold: **ADF Mapping Data Flows** standardize/clean/type-cast to **Silver Delta** with Alter Row upserts on business keys; **Gold views (joins, aggregations, dense-rank, Top-N)** optimized for **BI**.
- Centralized control via a **parent pipeline** chaining on-prem/API/SQL paths, **parameterizing file lists & batch sizes**, and parallelizing safe steps; less manual coordination and **predictable backfills/replays**.
- Shipped **ARM publish artifacts** and runbooks for **repeatable CI/CD**, with guardrails on **schema-drift vs validation, partitioning options**, and Gold snapshot overwrite strategy.

DBT Databricks Scd2-Data Quality: ([GitHub Link](#))

- Built a **Medallion** model with **dbt Core** on Databricks using env-driven **profiles.yml** and per-layer materializations.
- Enforced data quality with **dbt tests**: generic, **singular**, and **custom generic**.
- Implemented **SCD2** via **dbt snapshots** (timestamp strategy) for **time-travel analysis**.
- Drove reuse & speed with **Jinja macros**, **seeds**, **incremental** models, node selection, and **dbt build + docs/lineage** for CI.